



◀

0.

1.

2.

3.

4.

5.

6.

7.

8.

9.

10.

11.

12.

13.

14.

15.

## CSCI 2400 - Grunwald - Computer Systems

Home / My courses / Fall 2018 / CSCI2400-Fa18 / 12 November - 18 November / Malloc Lab

# Malloc Lab

## CS2400 - Malloc Lab: Writing a Dynamic Storage Allocator (short version)

### Introduction

In this lab you will be writing a dynamic storage allocator for C programs, i.e., your own version of the `malloc`, `free` and `realloc` routines. You are encouraged to explore the design space creatively and implement an allocator that is correct, efficient and fast.

### Logistics

You may work in a group of up to two people. Any clarifications and revisions to the assignment will be posted on the course Moodle.

### Hand Out Instructions

Download the `malloclab-handout.tar` file from the Moodle assignment page.

Start by copying `malloclab-handout.tar` to a directory in which you plan to do your work. Then give the command:

```
tar xvf malloclab-handout.tar
```

This will cause a number of files to be unpacked into the directory. The only file you will be modifying and handing in is `mm.c`. The `mdriver.c` program is a driver program that allows you to evaluate the performance of your solution. Use the command `make` to generate the driver code and run it with the command `./mdriver -V`. (The `-V` flag displays helpful summary information.)

Looking at the file `mm.c` you'll notice a C structure `team` into which you should insert the requested identifying information about the one or two individuals comprising your programming team. **Do this right away so you don't forget.**

When you have completed the lab, you will hand in only one file (`mm.c`), which contains your solution.

### How to Work on the Lab

Your dynamic storage allocator will consist of the following four functions, which are declared in `mm.h` and defined in `mm.c`.

```
int mm_init(void);
void *mm_malloc(size_t size);
void mm_free(void *ptr);
void *mm_realloc(void *ptr, size_t size);
```

The `mm.c` file we have given you implements the simplest but still functionally correct malloc package that we could think of. Using this as a starting place, modify these functions (and possibly define other private `static` functions), so that they obey the following semantics:

- **`mm_init`:** Before calling `mm_malloc`, `mm_realloc` or `mm_free`, the application program (i.e., the trace-driven driver program that you will use to evaluate your implementation) calls `mm_init` to perform any necessary initializations, such as allocating the initial heap area. The return value should be -1 if there was a problem in performing the initialization, 0 otherwise.
- **`mm_malloc`:** The `mm_malloc` routine returns a pointer to an allocated block payload of at least `size` bytes. The entire allocated block should lie within the heap region and should not overlap with any other allocated chunk. We will comparing your implementation to the version of `malloc` supplied in the standard C library (`libc`). Since the `libc` malloc always returns payload pointers that are aligned to 8 bytes, your malloc implementation should do likewise and always return 8-byte aligned pointers.
- **`mm_free`:** The `mm_free` routine frees the block pointed to by `ptr`. It returns nothing. This routine is only guaranteed to work when the passed pointer (`ptr`) was returned by an earlier call to `mm_malloc` or `mm_realloc` and has not yet been freed.
- **`mm_realloc`:** The `mm_realloc` routine returns a pointer to an allocated region of at least `size` bytes with the following constraints.
  - if `ptr` is NULL, the call is equivalent to `mm_malloc(size)`;
  - if `size` is equal to zero, the call is equivalent to `mm_free(ptr)`;
  - if `ptr` is not NULL, it must have been returned by an earlier call to `mm_malloc` or `mm_realloc`. The call to `mm_realloc` changes the size of the memory block pointed to by `ptr` (the item old block) to `size` bytes and returns the address of the new block. Notice that the address of the new block might be the same as the old block, or it might be different, depending on your implementation, the amount of internal fragmentation in the old block, and the size of the `realloc` request. The contents of the new block are the same as those of the old `ptr` block, up to the minimum of the old and new sizes. Everything else is uninitialized. For example, if the old block is 8 bytes and the new block is 12 bytes, then the first 8 bytes of the new block are identical to the first 8 bytes of the old block and the last 4 bytes are uninitialized. Similarly, if the old block is 8 bytes and the new block is 4 bytes, then the contents of the new block are identical to the first 4 bytes of the old block.

These semantics match the the semantics of the corresponding `libc` `malloc`, `realloc`, and `free` routines. Type `man malloc` to the shell for complete documentation or just google `malloc`.

### Heap Consistency Checker

Dynamic memory allocators are notoriously tricky beasts to program correctly and efficiently. They are difficult to program correctly because they involve a lot of unfused pointer manipulation. You will find it very helpful to write a heap checker that scans the heap and checks it for consistency.

Some examples of what a heap checker might check are:

- Is every block in the free list marked as free?
- Are there any contiguous free blocks that somehow escaped coalescing?
- Is every free block actually in the free list?
- Do the pointers in the free list point to valid free blocks?
- Do any allocated blocks overlap?
- Do the pointers in a heap block point to valid heap addresses?

Your heap checker will consist of the function `int mm_check(void)` in `mm.c`. It will check any invariants or consistency conditions you consider prudent. It returns a nonzero value if and only if your heap is consistent. You are not limited to the listed suggestions nor are you required to check all of them. You are encouraged to print out error messages when `mm_check` fails.

This consistency checker is for your own debugging during development. When you submit `mm.c`, make sure to remove any calls to `mm_check` as they will slow down your throughput.

## Support Routines

The memlib.c package simulates the memory system for your dynamic memory allocator. You can invoke the following functions in `memlib.c`:

- `void *mem_sbrk(int incr);`: Expands the heap by `incr` bytes, where `incr` is a positive non-zero integer and returns a generic pointer to the first byte of the newly allocated heap area. The semantics are identical to the Unix `sbrk` function, except that `mem_sbrk` accepts only a positive non-zero integer argument.
- `void *mem_heap_lo(void);`: Returns a generic pointer to the first byte in the heap.
- `void *mem_heap_hi(void);`: Returns a generic pointer to the last byte in the heap.
- `size_t mem_heapsize(void);`: Returns the current size of the heap in bytes.
- `size_t mem_pagesize(void);`: Returns the system's page size in bytes (4K on Linux systems).

## The Trace-driven Driver Program

The driver program `mdriver.c` in the `malloclabs-handout.tar` distribution tests your `mm.c` package for correctness, space utilization, and throughput. The driver program is controlled by a set of *trace files* that are included in the `malloclabs-handout.tar` distribution. Each trace file contains a sequence of allocate, reallocate, and free directions that instruct the driver to call your `mm_malloc`, `mm_realloc`, and `mm_free` routines in some sequence. The driver and the trace files are the same ones we will use when we grade your handin `mm.c` file.

The driver `mdriver.c` accepts the following command line arguments:

- `-t <tracedir>`: Look for the default trace files in directory `tracedir` instead of the default directory defined in `config.h`.
- `-f <tracefile>`: Use one particular `tracefile` for testing instead of the default set of tracefiles.
- `-h`: Print a summary of the command line arguments.
- `-l`: Run and measure `libc` malloc in addition to the student's malloc package.
- `-v`: Verbose output. Print a performance breakdown for each tracefile in a compact table.
- `-vv`: More verbose output. Prints additional diagnostic information as each trace file is processed. Useful during debugging for determining which trace file is causing your malloc package to fail.

## Programming Rules

- You should not change any of the interfaces in `mm.c`.
- You should not invoke any memory-management related library calls or system calls. This excludes the use of `malloc`, `calloc`, `free`, `realloc`, `sbrk`, `brk` or any variants of these calls in your code.
- For consistency with the `libc` `malloc` package, which returns blocks aligned on 8-byte boundaries, *your allocator must always return pointers that are aligned to 8-byte boundaries*. The driver will enforce this requirement for you.

## Evaluation

You'll be evaluated by having a functioning `malloc`.

The driver program summarizes the performance of your allocator by computing a *performance index*,  $P$ , which is a weighted sum of the space utilization and throughput

$$P = wU + (1 - w)\min(1, \frac{T}{T_{libc}})$$

where  $U$  is your space utilization,  $T$  is your throughput, and  $T_{libc}$  is the estimated throughput of `libc` malloc on your system on the default traces. The value for  $T_{libc}$  is a constant in the driver (600 Kops/s) that we established when we configured the program. Since we're using so many different machines, you should take this as a nominal" throughput for malloc on modern-day machines. The performance index favors space utilization over throughput. default of  $w=0.6w = 0.6$ .

Observing that both memory and CPU cycles are expensive system resources, we adopt this formula to encourage balanced optimization of both memory utilization and throughput. Ideally, the performance index will reach  $P = w + (1 - w) = 1$  or ( 100% ). Since each metric will contribute at most  $wu$  and  $1 - w$  to the performance index, respectively, you should not go to extremes to optimize either the memory utilization or the throughput only. To receive a good score, you must achieve a balance between utilization and throughput.

There is a scoring program, called `./RUN-MM` that will compile your program and run the test cases. This will report an *average score*. Your grade on the "does it work" portion of the machine problem is computed by the RUN-MM script using the average score.

The scoring function is based on using the <https://coding.csel.io> machines and specific implementations as goals:

- Score of 72, or ~65% is the implementation is the book, fleshed out and implemented. This is a first fit" implicit list allocator.
- Score of 80, or ~90%, is a next fit" implicit list allocator
- Score of 84, or ~100%, is an explicit" list allocator (see 9.9.13)
- Score of ~86, or ~105%, using circular" lists
- Score of 95, or 110%, is a tree based" allocator

You can earn up to 10% extra credit based on the score (i.e. a maximum score of 110).

# Handin Instructions

You will handin a zip file containing a valid mm.c, Makefile and other files needed to automatically build your program. You labs will be run by the TA's on the perf machine and grades recorded using the grading script.

*You must indicate the origin of any code you use or borrow from other people or sources. Failure to properly attribute the origin of code you retrieve from any source is grounds for receiving a zero.*

## Hints

- Use the `mdriver -f` option. During initial development, using tiny trace files will simplify debugging and testing. We have included two such trace files (`short1,2-bal.rep`) that you can use for initial debugging.
- Use the `mdriver -v` and `-vv` options. The `-v` option will give you a detailed summary for each trace file. The `-vv` will also indicate when each trace file is read, which will help you isolate errors.
- Compile with `gcc -g` and use a debugger. A debugger will help you isolate and identify out of bounds memory references.
- Understand every line of the malloc implementation in the textbook. The textbook has a detailed example of a simple allocator based on an implicit free list. Use this is a point of departure. Don't start working on your allocator until you understand everything about the simple implicit list allocator. That's good for 70% on this assignment.
- Encapsulate your pointer arithmetic in C preprocessor macros or inline functions. Pointer arithmetic in memory managers is confusing and error-prone because of all the casting that is necessary. You can reduce the complexity significantly by writing macros or, better yet, inline functions for your pointer operations. See the text for examples and look at the handout for provided samples.
- Do your implementation in stages. The first 9 traces contain requests to `malloc` and `free`. The last 2 traces contain requests for `realloc`, `malloc`, and `free`. We recommend that you start by getting your `malloc` and `free` routines working correctly and efficiently on the first 9 traces. Only then should you turn your attention to the `realloc` implementation. For starters, we've built a `realloc` on top of your existing `malloc` and `free` implementations. But to get really good performance, you will need to build a stand-alone `realloc`.
- Start early It is possible to write an efficient malloc package with a few pages of code. However, we can guarantee that it will be some of the most difficult and sophisticated code you have written so far in your career. So start early, and good luck!



## Submission status

Submission status	No attempt
Grading status	Not graded
Due date	Sunday, 9 December 2018, 11:59 PM
Time remaining	Assignment is overdue by: 6 days 22 hours
Last modified	-
Submission comments	<a href="#">Comments (0)</a>

## Follow Us



help@cs.colorado.edu

[Data retention summary](#)  
[Get the mobile app](#)